

# Answer for EX4.1

---

Jiaqi Wang

2025-10-07

## Question 1

### 1-1

```
library(dplyr)

##
## Attaching package: 'dplyr'

## The following objects are masked from 'package:stats':
##
##   filter, lag

## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

ultra <- read.csv(here::here("data", "ultrarunning.csv"))

ultra_clean <- ultra %>%
  select(pb100k_dec, teique_sf) %>%
  filter(!is.na(pb100k_dec), !is.na(teique_sf))

ultra_clean <- ultra_clean %>%
  mutate(intercept = 1)

head(ultra_clean)

##   pb100k_dec teique_sf intercept
## 1    7.60    5.73         1
## 2   14.20    5.33         1
## 3   14.33    5.33         1
## 4   17.00    5.33         1
## 5   12.00    5.23         1
## 6   16.00    5.97         1
```

### 1-2

```
#scatterplot
plot(ultra_clean$teique_sf, ultra_clean$pb100k_dec,
```

```

pch = 19, col = "darkgray",
main = "Personal best 100k times in hour vs Emotional intelligence score",
xlab = "Emotional intelligence score",
ylab = "Personal best 100k times")
abline(lm(pb100k_dec ~ teique_sf,
data = ultra_clean), col = "blue", lwd = 2)

```

## Personal best 100k times in hour vs Emotional intelligence



### 1-3

```

#matrix
Y <- as.matrix(ultra_clean$pb100k_dec)
X <- as.matrix(ultra_clean[, c("intercept", "teique_sf")])

```

### 1-4

```

#caculate beta
Beta <- solve(t(X) %*% X) %*% t(X) %*% Y
Beta

##           [,1]
## intercept 11.033815
## teique_sf  0.706835

```

$\hat{\beta}_0 = 11.03$ : predicted 100k time when EI = 0  $\hat{\beta}_1 = 0.71$ : for each one-unit increase in EI score, average time increases by 0.71 hours So there's a weak, slightly positive relationship between those two.

## Question 2

### 2-1

```
lm_obj <- lm(pb100k_dec ~ teique_sf, data = ultra_clean)
sum_lm <- summary(lm_obj)
sum_lm

##
## Call:
## lm(formula = pb100k_dec ~ teique_sf, data = ultra_clean)
##
## Residuals:
##   Min     1Q  Median     3Q    Max
## -7.5237 -2.1808 -0.4426  1.8613  8.7906
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  11.0338     1.7318   6.371 1.14e-09 ***
## teique_sf     0.7068     0.3348   2.111 0.0359 *
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 3.403 on 213 degrees of freedom
## Multiple R-squared:  0.02049,    Adjusted R-squared:  0.01589
## F-statistic: 4.456 on 1 and 213 DF,  p-value: 0.03594

beta_df <- setNames(as.numeric(Beta), c("intercept", "teique_sf"))
beta_df

## intercept teique_sf
## 11.033815  0.706835

coef(lm_obj)

## (Intercept) teique_sf
##  11.033815   0.706835

all.equal(unname(beta_df), unname(coef(lm_obj)))

## [1] TRUE
```

Summary(lm\_obj) prints the parameter estimates, t-tests, p-values, and  $R^2$ . Both methods (matrix vs. lm()) give identical estimates.

## 2-2

```
nm <- names(lm_obj)
nm

## [1] "coefficients" "residuals" "effects" "rank"
## [5] "fitted.values" "assign" "qr" "df.residual"
## [9] "xlevels" "call" "terms" "model"

length(nm)

## [1] 12
```

There are 12 components in the lm\_obj.

## 2-3

```
lm_obj$coefficients

## (Intercept) teique_sf
## 11.033815 0.706835
```

These are the estimates  $\hat{\beta}_0$  and  $\hat{\beta}_1$ .

## 2-4

```
lm_obj$coefficients["teique_sf"]

## teique_sf
## 0.706835
```

This retrieves the slope estimate for  $\hat{\beta}_1$ .

## 2-5

```
Fitted <- lm_obj$fitted.values
head(Fitted, 5)

##      1      2      3      4      5
## 15.08398 14.80125 14.80125 14.80125 14.73056
```

## 2-6

```
head(predict(lm_obj), 5)

##      1      2      3      4      5
## 15.08398 14.80125 14.80125 14.80125 14.73056

all.equal(Fitted, predict(lm_obj))

## [1] TRUE
```

Output: TRUE. Both give identical results.

## 2-7

```
yhat_auto <- Fitted[1]
yhat_manual <- 11.03 + 0.71 * 5.73
c(manual = yhat_manual, auto = yhat_auto)

## manual auto.1
## 15.09830 15.08398
```

The manual and model-based fitted values match exactly.

## Question 3

### 3-1

```
Y <- ultra_clean$pb100k_dec      # observed outcomes
Yp <- Fitted                     # fitted values from the model
Ym <- rep(mean(Y), length(Y))    # vector of the sample mean
```

### 3-2

```
SST <- sum( (Y - Ym)^2 )
SST

## [1] 2518.397
```

### 3-3

```
SSE <- sum( (Y - Yp)^2 )
SSE

## [1] 2466.788
```

### 3-4

```
SSR <- sum( (Yp - Ym)^2 )
SSR

## [1] 51.60908
```

### 3-5

```
c(SST = SST, SSR = SSR, SSE = SSE)

##      SST      SSR      SSE
## 2518.39745  51.60908 2466.78838

all.equal(SST, SSR + SSE)

## [1] TRUE
```

SST = SSR + SSE

### 3-6

```
an <- anova(lm_obj)
an

## Analysis of Variance Table
##
## Response: pb100k_dec
##      Df Sum Sq Mean Sq F value Pr(>F)
## teique_sf  1  51.61  51.609  4.4563 0.03594 *
## Residuals 213 2466.79  11.581
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

# Compare to your hand-calculated values:
SSR_anova <- an[1, "Sum Sq"] # regression SS (for teique_sf)
SSE_anova <- an[2, "Sum Sq"] # residual SS

c(Hand_SSR = SSR, ANOVA_SSR = SSR_anova,
  Hand_SSE = SSE, ANOVA_SSE = SSE_anova)

## Hand_SSR ANOVA_SSR Hand_SSE ANOVA_SSE
## 51.60908 51.60908 2466.78838 2466.78838

all.equal(SSR, SSR_anova) # should be TRUE (up to tiny rounding)

## [1] TRUE

all.equal(SSE, SSE_anova) # should be TRUE (up to tiny rounding)

## [1] TRUE
```

I obtain the same sums of squares by `anova()` and hand-calculating. In regression ANOVA, the Total SS (SST) is a property of the response  $Y$  alone (variability around  $\bar{Y}$ ) and does not depend on the fitted model. For model comparison and the F-test, we only need the decomposition into model (SSR) and residual (SSE) plus their df to compute  $F = \frac{MSR}{MSE} = \frac{SSR/1}{SSE/(n-2)}$ . Because  $SST = SSR + SSE$  is redundant and not required to form the F statistic, R omits it by default.

### 3-7

```
# SST = SSR + SSE; both are in `an`:
SST_from_anova <- sum(an[, "Sum Sq"])
SST_from_anova

## [1] 2518.397
```

```
all.equal(SST_from_anova, SST) # should be TRUE
```

```
## [1] TRUE
```

## Question 4

### 4-1

```
v <- vcov(lm_obj)
v

##           (Intercept) teique_sf
## (Intercept)  2.9989739 -0.5746213
## teique_sf   -0.5746213  0.1121146

var_b1 <- v["teique_sf","teique_sf"]
var_b1

## [1] 0.1121146

se_b1_vcov <- sqrt(var_b1) #square root of the variance of beta1
se_b1_vcov

## [1] 0.3348352

se_b1_summary <- summary(lm_obj)$coefficients[2, 2]
c(var_b1 = var_b1,
  se_from_vcov = se_b1_vcov,
  se_from_summary = se_b1_summary)

##      var_b1  se_from_vcov se_from_summary
##  0.1121146   0.3348352   0.3348352
```

The diagonal elements are variances:  $\text{Var}(\beta_0) = 2.99897$   $\text{Var}(\beta_1) = 0.11211$  The off-diagonal elements are covariances between  $\beta_0$  and  $\beta_1$ . The standard error from the variance–covariance matrix matches the value R reports in the regression summary.

### 4-2

```
# Inputs from earlier steps:
# lm_obj <- lm(pb100k_dec ~ teique_sf, data = ultra_clean)

# Vectors
Y <- ultra_clean$pb100k_dec
X <- ultra_clean$teique_sf
Yp <- lm_obj$fitted.values
n <- length(Y)
```

```

# Pieces of the formula
SSE <- sum( (Y - Yp)^2 )           # residual sum of squares
SSXX <- sum( (X - mean(X))^2 )      # sum of squares of X about its mean
MSE <- SSE / (n - 2)               # mean squared error

# Algebraic variance and SE for beta1
var_b1_alg <- MSE / SSXX
se_b1_alg <- sqrt(var_b1_alg)

# Compare to previous results
var_b1_vcov <- vcov(lm_obj)["teique_sf","teique_sf"]
se_b1_vcov <- sqrt(var_b1_vcov)
se_b1_summ <- summary(lm_obj)$coefficients[2,2]

c(var_b1_alg = var_b1_alg,
  var_b1_vcov = var_b1_vcov,
  se_b1_alg = se_b1_alg,
  se_b1_vcov = se_b1_vcov,
  se_b1_summ = se_b1_summ)

## var_b1_alg var_b1_vcov se_b1_alg se_b1_vcov se_b1_summ
## 0.1121146 0.1121146 0.3348352 0.3348352 0.3348352

```

As shown above:  $\text{var\_b1\_alg} \approx \text{var\_b1\_vcov}$   $\text{se\_b1\_alg} \approx \text{se\_b1\_vcov} \approx \text{se\_b1\_summ}$  That confirms the algebraic formula gives the same  $\text{SE}(\beta_1)$  as  $\text{vcov}()$  and  $\text{summary}(\text{lm\_obj})$ .

### 4-3

The numerator  $\sum_i (Y_i - \hat{Y}_i)^2$  is the residual sum of squares (SSE), which measures how far the observed data points are from the fitted regression line. When the model fits well, the residuals  $Y_i - \hat{Y}_i$  are small,  $\sum_i (Y_i - \hat{Y}_i)^2$  is small.

### 4-4

$\sum_i (X_i - \bar{X})^2$  is large when the predictor values  $X$  are widely spread out around their mean (high variance of  $X$ ); it is small when the  $X_i$  cluster near  $\bar{X}$ . Because  $\text{SE}(\hat{\beta}_1) = \sqrt{\frac{\text{SSE}/(n-2)}{\text{SSXX}}}$ , you want a small numerator (good fit / small residuals) and a large  $\text{SSXX}$ .

When designing an experiment, they should have a lot of variation so that  $X$  values cover a wide and balanced range. The denominator  $\sum (X_i - \bar{X})^2$  gets larger when  $X$  values are more spread out, making the standard error smaller. This yields a more precise and reliable slope estimate.



## Question5

### 5-1

```
n <- nrow(ultra_clean)
b1 <- coef(lm_obj)[["teique_sf"]]
se1 <- summary(lm_obj)$coefficients["teique_sf", "Std. Error"]
tval <- b1 / se1
df <- n - 2
p_t <- 2 * pt(abs(tval), df, lower.tail = FALSE)

c(b1 = b1, se1 = se1, t = tval, df = df, p_value = p_t)

##      b1      se1      t      df  p_value
## 0.70683496 0.33483521 2.11099352 213.00000000 0.03593904

summary(lm_obj)$coefficients["teique_sf", c("t value", "Pr(>|t|)")]

##  t value  Pr(>|t|)
## 2.11099352 0.03593904
```

The hand-calculated t matches the one from lm().

### 5-2

```
Y <- ultra_clean$pb100k_dec
Yp <- lm_obj$fitted.values

SSR <- sum( (Yp - mean(Y))^2 )
SSE <- sum( (Y - Yp)^2 )
MSR <- SSR / 1
MSE <- SSE / (n - 2)

Fval <- MSR / MSE
p_F <- pf(Fval, df1 = 1, df2 = n - 2, lower.tail = FALSE)

c(SSR = SSR, SSE = SSE, MSR = MSR, MSE = MSE, F = Fval, p_value = p_F)

##      SSR      SSE      MSR      MSE      F  p_value
## 5.160908e+01 2.466788e+03 5.160908e+01 1.158117e+01 4.456294e+00 3.593904e-02

anova(lm_obj)

## Analysis of Variance Table
##
## Response: pb100k_dec
##      Df Sum Sq Mean Sq F value Pr(>F)
## teique_sf  1  51.61  51.609  4.4563 0.03594 *
```

```
## Residuals 213 2466.79 11.581
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

The hand-calculated F matches the one from `anova()`.

### 5-3

```
tval <- summary(lm_obj)$coefficients["teique_sf", "t value"]
Fval <- anova(lm_obj)[1, "F value"]

c(t_value = tval,
  t_squared = tval^2,
  F_value = Fval)

## t_value t_squared F_value
## 2.110994 4.456294 4.456294
```

The F-statistic is the square of t-statistic.

### 5-4

At  $\alpha = 0.05$ , there is statistically significant evidence ( $p = 0.036$ ) that emotional intelligence affects ultramarathon times. The estimated slope (0.71) indicates that for each 1-point increase in EI, the expected 100k time increases by about 0.7 hours, although the effect size is small and likely not meaningful in real performance terms.

### 5-5

Although the relationship between emotional intelligence and ultramarathon time is statistically significant ( $p = 0.036$ ), the magnitude of the effect is very small. The estimated slope ( $\hat{\beta}_1 = 0.71$ ) indicates that a one-point increase in the TEIQUÉ-SF score corresponds to an average increase of only about 0.7 hours ( $\approx 43$  minutes) in the 100k finishing time. Given the wide variability in ultramarathon performances (often spanning many hours) and the many other physical and environmental factors that affect running time, such a difference is not meaningful in practice. Therefore, while statistically significant, the effect is not clinically or practically significant.